

## Model Development Phase Template

Date	18 July 2024
Team ID	740111
Project Title	Unvieling Airbnb Price Patterns : Machine Learning Models For Forecasting
Maximum Marks	4 Marks

### Initial Model Training Code, Model Validation and Evaluation Report

The initial model training code will be showcased in the future through a screenshot. The model validation and evaluation report will include classification reports, accuracy, and confusion matrices for multiple models, presented through respective screenshots.

#### Initial Model Training Code:

Linear regression:

```
import numpy as np
from sklearn.linear_model import LinearRegression
from sklearn import metrics

# Assuming x_train_scaled, y_train, x_test_scaled, and y_test are already defined

# Initialize the Linear Regression model
lr = LinearRegression()

# Fit the model to the training data
lr.fit(x_train_scaled, y_train)

# Predict using the test data
y_pred_lr = lr.predict(x_test_scaled)

# Calculate evaluation metrics
mae_lr = metrics.mean_absolute_error(y_test, y_pred_lr)
mse_lr = metrics.mean_squared_error(y_test, y_pred_lr)
rmse_lr = np.sqrt(mse_lr) # RMSE is already calculated correctly
r2_lr = metrics.r2_score(y_test, y_pred_lr)

# Print evaluation metrics
print('\nMean Absolute Error of Linear Regression:', mae_lr)
print('\nMean Squared Error of Linear Regression:', mse_lr)
print('\nRoot Mean Squared Error of Linear Regression:', rmse_lr)
print('\nR2 Score of Linear Regression:', r2_lr)
```

#### Random Forest Regressor

```
rf=RandomForestRegressor()
rf.fit(x_train_scaled,y_train)
y_pred_rf= rf.predict(x_test_scaled)
mae_rf= metrics.mean_absolute_error(y_test, y_pred_rf)
mse_rf= metrics.mean_squared_error(y_test, y_pred_rf)
rmse_rf =np.sqrt(metrics.mean_squared_error(y_test, y_pred_rf))
r2_rf= metrics.r2_score(y_test, y_pred_rf)
print('\nMean Absolute Error of Random Forest Regressor:',mae_rf)
print('\nMean Squarred Error of Random Forest Regressor :', mse_rf)
print('\nRoot Mean Squarred Error of Random Forest Regressor:', rmse_rf)
print('\nR2 Score of Random Forest Regressor :',r2_rf)
```

## Polynomial Regression:

```
from sklearn.linear_model import Ridge
model=Pipeline([
    ('poly', PolynomialFeatures()),
    ('ridge', Ridge(fit_intercept=True))
])
param_grid={
    'poly_degree': [1, 2, 3],
    'ridge_alpha': [0.1,0.5,1.0,2.0]
}
#Perform grid search with 5-fold cross-validation
poly_tuned=GridSearchCV(model, param_grid, cv=5)
#Training and Testing
poly_tuned.fit(x_train_scaled, y_train)
y_pred_poly= poly_tuned.predict(x_test_scaled)
mae_poly= metrics.mean_absolute_error(y_test, y_pred_poly)
mse_poly= metrics.mean_squared_error(y_test, y_pred_poly)
rmse_poly= np.sqrt(metrics.mean_squared_error(y_test, y_pred_poly))
r2_poly= metrics.r2_score(y_test, y_pred_poly)
print('\nMean Absolute Error of Polynomial Regression:',mae_poly)
print('\nMean Squared Error of Polynomial Regression:',mse_poly)
print('\nRoot Mean Squared Error of Polynomial Regression:',rmse_poly)
print('\nR2 Score of Polynomial Regression :',r2_poly)
```

## Gradient Boosting Regressor:

```
from sklearn.ensemble import GradientBoostingRegressor
from sklearn import metrics

# Assuming x_train_scaled, x_test_scaled, y_train, y_test are defined and properly scaled

# Initialize GradientBoostingRegressor
gb = GradientBoostingRegressor(n_estimators=200, learning_rate=0.2, max_depth=10)

# Fit the model
gb.fit(x_train_scaled, y_train)

# Predict using the model
y_pred_gb = gb.predict(x_test_scaled)

# Calculate evaluation metrics
mae_gb = metrics.mean_absolute_error(y_test, y_pred_gb)
mse_gb = metrics.mean_squared_error(y_test, y_pred_gb) # Corrected variable name
rmse_gb = np.sqrt(metrics.mean_squared_error(y_test, y_pred_gb))
r2_gb = metrics.r2_score(y_test, y_pred_gb)

# Print the results
print('\nMean Absolute Error of Gradient Boosting:', mae_gb)
print('\nMean Squared Error of Gradient Boosting:', mse_gb) # Corrected print statement
print('\nRoot Mean Squared Error of Gradient Boosting:', rmse_gb)
print('\nR2 Score of Gradient Boosting:', r2_gb)
```

## Catboost Regressor:

```
from catboost import CatBoostRegressor
from sklearn.model_selection import cross_val_score, KFold
from sklearn import metrics
import numpy as np
model_CBR = CatBoostRegressor()
model_CBR.fit(x_train_scaled, y_train)
cross_val_score(model_CBR, x_train_scaled, y_train,
                 scoring='r2',
                 cv=KFold(n_splits=5,
                          shuffle=True,
                          random_state=2022,
                          ))
y_pred_cbr = model_CBR.predict(x_test_scaled)
mae_cbr = metrics.mean_absolute_error(y_test, y_pred_cbr)
mse_cbr = metrics.mean_squared_error(y_test, y_pred_cbr)
rmse_cbr = np.sqrt(metrics.mean_squared_error(y_test, y_pred_cbr))
r2_cbr = metrics.r2_score(y_test, y_pred_cbr)
print('\nMean Absolute Error of CatBoost Regressor:', mae_cbr)
print('\nMean Squared Error of CatBoost Regressor:', mse_cbr)
print('\nRoot Mean Squared Error of CatBoost Regressor:', rmse_cbr)
print('\nR2 Score of CatBoost Regressor:', r2_cbr)
```

### XGBoost Regressor:

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn import metrics
from xgboost import XGBRegressor
xgb = XGBRegressor(objective='reg:squarederror')
xgb.fit(x_train_scaled, y_train)
y_pred_xgb = xgb.predict(x_test_scaled)
mae_xgb = metrics.mean_absolute_error(y_test, y_pred_xgb)
mse_xgb = metrics.mean_squared_error(y_test, y_pred_xgb)
rmse_xgb = np.sqrt(mse_xgb)
r2_xgb = metrics.r2_score(y_test, y_pred_xgb)
print('\nMean Absolute Error of XGBoost Regressor:', mae_xgb)
print('\nMean Squared Error of XGBoost Regressor:', mse_xgb)
print('\nRoot Mean Squared Error of XGBoost Regressor:', rmse_xgb)
print('\nR2 Score of XGBoost Regressor:', r2_xgb)
```

Model	Classification Report	F1 Score	Confusion Matrix
-------	-----------------------	----------	------------------

Linear regression	<pre> Mean Absolute Error of Linear Regression: 0.2496850406676611 Mean Squared Error of Linear Regression: 0.0831286214866343 Root Mean Squared Error of Linear Regression: 0.28832034525269684 R2 Score of Linear Regression: -0.00040043704639391997 </pre>	-	-
-------------------	--	---	---

### Model Validation and Evaluation Report:

Forest regression	<pre> Mean Absolute Error of Random Forest Regressor: 0.25042532597042355 Mean Squarred Error of Random Forest Regressor : 0.0840184970084306 Root Mean Squarred Error of Random Forest Regressor: 0.2898594435384685 R2 Score of Random Forest Regressor : -0.011109526707709039 </pre>	--	-
Polynomial regression	<pre> Mean Absolute Error of Polynomial Regression: 0.24968503845646894 Mean Squarred Error of Polynomial Regression: 0.08312861967771229 Root Mean Squarred Error of Polynomial Regression: 0.288320342115696 R2 Score of Polynomial Regression : -0.00040041527716039305 </pre>	-	-
Gradient boosting	<pre> Mean Absolute Error of Gradient Boosting: 0.2570986980685457 Mean Squared Error of Gradient Boosting: 0.0904058531652557 Root Mean Squared Error of Gradient Boosting: 0.30067566107893684 R2 Score of Gradient Boosting: -0.08797732237885714 </pre>	-	-

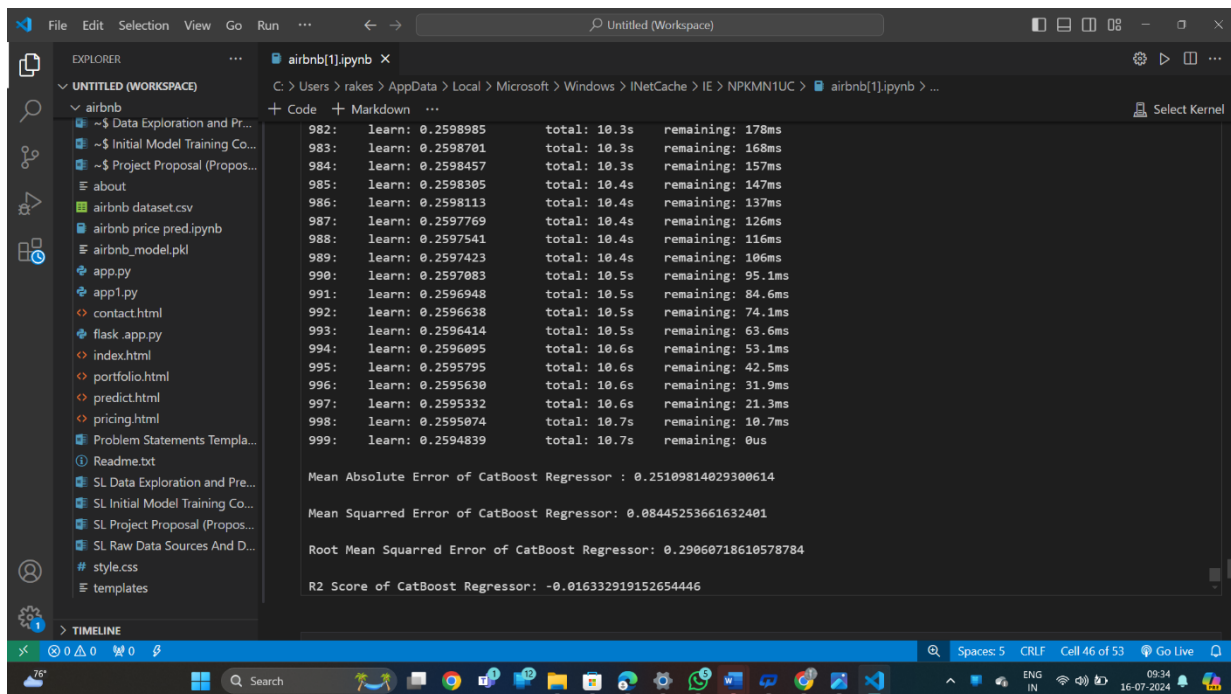
```
Mean Absolute Error of XGBoost Regressor: 0.2536936827472256

Mean Squared Error of XGBoost Regressor: 0.08700855364433946

Root Mean Squared Error of XGBoost Regressor: 0.2949721235037973

R2 Score of XGBoost Regressor: -0.04709296913538896
```

Catboost regressor:



The screenshot shows a Jupyter Notebook titled 'airbnb[1].ipynb' in a web browser. The notebook contains a table of performance metrics for a CatBoost regressor. The table has four columns: 'learn', 'total', 'remaining', and 'remaining'. The rows show the progress of the training process, with the final row indicating the completion of the training. Below the table, the final performance metrics are displayed.

learn	total	remaining
0.2598985	10.3s	178ms
0.2598701	10.3s	168ms
0.2598457	10.3s	157ms
0.2598305	10.4s	147ms
0.2598113	10.4s	137ms
0.2597769	10.4s	126ms
0.2597541	10.4s	116ms
0.2597423	10.4s	106ms
0.2597083	10.5s	95.1ms
0.2596948	10.5s	84.6ms
0.2596638	10.5s	74.1ms
0.2596414	10.5s	63.6ms
0.2596095	10.6s	53.1ms
0.2595795	10.6s	42.5ms
0.2595630	10.6s	31.9ms
0.2595332	10.6s	21.3ms
0.2595074	10.7s	10.7ms
0.2594839	10.7s	0us

Mean Absolute Error of CatBoost Regressor : 0.25109814029380614

Mean Squared Error of CatBoost Regressor: 0.08445253661632481

Root Mean Squared Error of CatBoost Regressor: 0.29060718610578784

R2 Score of CatBoost Regressor: -0.016332919152654446